

```
In [125]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import cm
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings("ignore")
```

```
In [126]: df = pd.read_csv('GEMIRAN0819.csv')
df
```

Out[126]:

	Unnamed: 0	yrsurv	age	gender	irstrata	fearfail	suskill	opport	knowent	discent	exre
0	3	2008	28	2	2	2	1	1	1	0	
1	20	2008	38	1	2	2	1	1	1	0	
2	25	2008	26	2	6	2	1	1	1	0	
3	30	2008	25	1	6	2	1	1	2	0	
4	35	2008	45	1	2	-1	2	-1	1	0	
...	...	...	...	...	...	...	...	...	...	...	...
12095	39411	2019	54	2	6	1	1	2	1	0	
12096	39414	2019	21	2	6	1	2	2	1	0	
12097	39416	2019	24	2	6	2	1	1	1	0	
12098	39429	2019	31	2	6	2	1	2	2	0	
12099	39436	2019	32	2	6	2	2	2	1	0	

12100 rows × 12 columns

```
In [127]: df=df.drop(['Unnamed: 0'],axis=1)
df
```

Out[127]:

	ysrurv	age	gender	irstrata	fearfail	suskill	opport	knowent	discent	exreason	bstart
0	2008	28	2	2	2	1	1	1	0	0	1
1	2008	38	1	2	2	1	1	1	0	0	1
2	2008	26	2	6	2	1	1	1	0	0	1
3	2008	25	1	6	2	1	1	2	0	0	2
4	2008	45	1	2	-1	2	-1	1	0	0	1
...	...	...	...	...	...	...	...	...	...	...	..
12095	2019	54	2	6	1	1	2	1	0	0	1
12096	2019	21	2	6	1	2	2	1	0	0	1
12097	2019	24	2	6	2	1	1	1	0	0	1
12098	2019	31	2	6	2	1	2	2	0	0	1
12099	2019	32	2	6	2	2	2	1	0	0	2

12100 rows × 11 columns


```
In [128]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12100 entries, 0 to 12099
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ysrurv      12100 non-null  int64
1   age         12100 non-null  int64
2   gender      12100 non-null  int64
3   irstrata    12100 non-null  int64
4   fearfail    12100 non-null  int64
5   suskill     12100 non-null  int64
6   opport      12100 non-null  int64
7   knowent     12100 non-null  int64
8   discent     12100 non-null  int64
9   exreason    12100 non-null  int64
10  bstart      12100 non-null  int64
dtypes: int64(11)
memory usage: 1.0 MB
```

```
In [129]: #Let's make our correlation matrix a little prettier
corr_matrix = df[['discent', 'exreason', 'fearfail', 'opport', 'knowent', 'suskill',
                  'bstart', 'istrata', 'yrsurv', 'gender', 'age']].corr(method='spearman')

fig, ax = plt.subplots(figsize=(8, 8))
ax = sns.heatmap(corr_matrix,
                  annot=True,
                  linewidths=0.5,
                  fmt=".2f",
                  cmap="YlGnBu");
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

Out[129]: (11.5, -0.5)



```
In [130]: istrata_table= pd.crosstab(index=df['istrata'], columns=df['discent' ])
#istrata_table['ratio']=istrata_table[1]/(istrata_table[0]+istrata_table
[1])
istrata_table.reset_index(inplace=True)
```

```
In [131]: instrata_table.style.background_gradient(cmap = 'Blues')
```

```
Out[131]:
```

	discent	instrata	0	1
0	1	1991	172	
1	2	2440	252	
2	3	2167	200	
3	4	1487	162	
4	5	1765	190	
5	6	1172	102	

```
In [132]: instrata_table= pd.crosstab(index=df['instrata'], columns=df['fearfail' ])
#instrata_table['ratio']=instrata_table[1]/(instrata_table[0]+instrata_table[1])
instrata_table.reset_index(inplace=True)
```

```
In [133]: instrata_table.style.background_gradient(cmap = 'Blues')
```

```
Out[133]:
```

	fearfail	instrata	-2	-1	0	1	2
0	1	4	33	4	811	1311	
1	2	23	60	33	921	1655	
2	3	19	155	2	661	1530	
3	4	9	110	7	512	1011	
4	5	12	70	8	569	1296	
5	6	8	28	2	382	854	

```
In [134]: instrata_table= pd.crosstab(index=df['instrata'], columns=df['bstart'])
#instrata_table['ratio']=instrata_table[1]/(instrata_table[0]+instrata_table[1])
instrata_table.reset_index(inplace=True)
```

```
In [135]: instrata_table.style.background_gradient(cmap = 'Blues')
```

```
Out[135]:
```

	bstart	instrata	-2	-1	1	2
0	1	4	4	944	1211	
1	2	4	3	1261	1424	
2	3	1	6	1055	1305	
3	4	0	3	737	909	
4	5	6	13	930	1006	
5	6	3	6	589	676	

```
In [136]: instrata_table= pd.crosstab(index=df['instrata'], columns=df['knowent'], margins_name="Total")
#instrata_table['ratio']=instrata_table[1]/(instrata_table[0]+instrata_table[1])
instrata_table.reset_index(inplace=True)
```

```
In [137]: instrata_table.style.background_gradient(cmap = 'Blues')
```

Out[137]:

	knowent	instrata	-2	-1	0	1	2
0	1	5	15	0	1097	1046	
1	2	12	45	2	1497	1136	
2	3	14	54	0	1162	1137	
3	4	6	28	1	943	671	
4	5	13	39	1	1134	768	
5	6	5	25	2	792	450	

```
In [138]: pd.crosstab(df.instrata, df['knowent'], margins=True, margins_name="Total").style.background_gradient(cmap = 'Blues')
```

Out[138]:

knowent	-2	-1	0	1	2	Total
instrata						
1	5	15	0	1097	1046	2163
2	12	45	2	1497	1136	2692
3	14	54	0	1162	1137	2367
4	6	28	1	943	671	1649
5	13	39	1	1134	768	1955
6	5	25	2	792	450	1274
Total	55	206	6	6625	5208	12100

```
In [139]: instrata_table= pd.crosstab(index=df['instrata'], columns=df['suskill'])
#instrata_table['ratio']=instrata_table[1]/(instrata_table[0]+instrata_table[1])
instrata_table.reset_index(inplace=True)
```

```
In [140]: instrata_table.style.background_gradient(cmap = 'Blues')
```

```
Out[140]:
```

suskill	instrata	-2	-1	0	1	2
0	1	9	19	9	1496	630
1	2	22	52	30	2001	587
2	3	21	245	6	1549	546
3	4	5	109	2	1204	329
4	5	11	60	10	1355	519
5	6	26	51	2	906	289

```
In [141]: instrata_table= pd.crosstab(index=df['instrata'], columns=df['opport' ])
#instrata_table['ratio']=instrata_table[1]/(instrata_table[0]+instrata_table
[1])*100
instrata_table.reset_index(inplace=True)
```

```
In [142]: instrata_table.style.background_gradient(cmap = 'Blues')
```

```
Out[142]:
```

opport	instrata	-2	-1	0	1	2
0	1	11	110	7	783	1252
1	2	18	254	43	1002	1375
2	3	41	409	7	824	1086
3	4	6	212	7	629	795
4	5	15	161	12	727	1040
5	6	19	183	3	465	604

```
In [143]: instrata_table= pd.crosstab(index=df['instrata'], columns=df['gender' ])
#instrata_table['ratio']=instrata_table[1]/(instrata_table[0]+instrata_table
[1])*100
instrata_table.reset_index(inplace=True)
```

```
In [144]: pd.crosstab(df.opport, df['discent'], margins=True, margins_name="Total").styl
e.background_gradient(cmap = 'Blues')
```

```
Out[144]:
```

discent	0	1	Total
opport			
-2	106	4	110
-1	1264	65	1329
0	71	8	79
1	3982	448	4430
2	5599	553	6152
Total	11022	1078	12100

```
In [145]: pd.crosstab(df.gender, df['exreason'], margins=True, margins_name="Total").style.background_gradient(cmap='Blues')
```

Out[145]:

exreason	0	1	2	3	4	5	6	7	8	9	10	Total
gender												
1	8398	13	110	7	29	35	141	37	354	61	17	9202
2	2683	5	50	2	2	8	41	11	70	24	2	2898
Total	11081	18	160	9	31	43	182	48	424	85	19	12100

```
In [146]: pd.crosstab(df.exreason, df['discent']).style.background_gradient(cmap='Blues')
```

Out[146]:

discent	0	1
exreason		
0	11019	62
1	0	18
2	1	159
3	0	9
4	0	31
5	0	43
6	0	182
7	0	48
8	1	423
9	0	85
10	1	18

```
In [147]: yrsurv_table= pd.crosstab(index=df['yrsurv'], columns=df['discent' ])
#yrsurv_table['ratio']=yrsurv_table[1]/(yrsurv_table[0]+yrsurv_table[1])
yrsurv_table.reset_index(inplace=True)
```

```
In [148]: yrsurv_table.style.background_gradient(cmap = 'Blues')
```

```
Out[148]:
```

	discent	yrsurv	0	1
0	2008	621	56	
1	2009	761	76	
2	2010	1054	88	
3	2011	1239	92	
4	2012	687	56	
5	2013	969	87	
6	2014	860	88	
7	2015	1036	116	
8	2016	1202	140	
9	2017	998	111	
10	2018	877	77	
11	2019	718	91	

```
In [149]: yrsurv_table= pd.crosstab(index=df['yrsurv'], columns=df['exreason' ])
#yrsurv_table['ratio']=yrsurv_table[1]/(yrsurv_table[0]+yrsurv_table[1])
yrsurv_table.reset_index(inplace=True)
```

```
In [150]: yrsurv_table.style.background_gradient(cmap = 'Blues')
```

```
Out[150]:
```

	exreason	yrsurv	0	1	2	3	4	5	6	7	8	9	10
0	2008	622	0	6	1	4	4	10	1	19	8	2	
1	2009	779	0	10	0	4	2	12	1	17	10	2	
2	2010	1058	0	11	0	5	8	21	1	32	5	1	
3	2011	1264	0	17	4	7	4	4	4	21	5	1	
4	2012	688	0	11	0	0	2	9	6	19	8	0	
5	2013	970	0	7	0	1	0	22	9	39	5	3	
6	2014	860	0	17	0	1	0	15	4	45	4	2	
7	2015	1039	3	21	1	3	3	18	4	52	7	1	
8	2016	1203	3	25	1	3	8	16	5	63	10	5	
9	2017	998	4	13	0	0	6	19	8	52	8	1	
10	2018	878	0	10	0	0	4	22	2	28	9	1	
11	2019	722	8	12	2	3	2	14	3	37	6	0	

```
In [151]: category = pd.cut(df.age,
                             bins=[18,25,35,45,55,65,99],
                             labels=['18-24', '25-34', '35-44', '45-54', '55-64', '65+'])
df.insert(5, 'Age-Group', category)
```



In [152]: df


Out[152]:

	yrsurv	age	gender	irstrata	fearfail	Age-Group	suskill	opport	knowent	discent	exreasor
0	2008	28	2	2	2	25-34	1	1	1	0	(
1	2008	38	1	2	2	35-44	1	1	1	0	(
2	2008	26	2	6	2	25-34	1	1	1	0	(
3	2008	25	1	6	2	18-24	1	1	2	0	(
4	2008	45	1	2	-1	35-44	2	-1	1	0	(
...	...	...	...	...	...	...	...	...	...	...	..
12095	2019	54	2	6	1	45-54	1	2	1	0	(
12096	2019	21	2	6	1	18-24	2	2	1	0	(
12097	2019	24	2	6	2	18-24	1	1	1	0	(
12098	2019	31	2	6	2	25-34	1	2	2	0	(
12099	2019	32	2	6	2	25-34	2	2	1	0	(

12100 rows × 12 columns

In [153]: pd.crosstab(df.gender, df['Age-Group'], margins=False).plot(kind='bar')

Out[153]: &lt;matplotlib.axes.\_subplots.AxesSubplot at 0x14c887ad7c8&gt;




```
In [156]: import matplotlib.path as mpath
q=df.groupby('yrsurv').agg({'discent':'sum'})

star = mpath.Path.unit_regular_star(6)
circle = mpath.Path.unit_circle()
# concatenate the circle with an internal cutout of the star
verts = np.concatenate([circle.vertices, star.vertices[:::-1, ...]])
codes = np.concatenate([circle.codes, star.codes])
cut_star = mpath.Path(verts, codes)

figsize = (8,8)
plt.plot(q, '--r', marker=cut_star, markersize=15,C='#d62728')


plt.ylabel('Number of Exit')
plt.title('Exit trends by year')
plt.show()
```



```
In [119]: fig = plt.subplots(figsize=(10,5))

sns.countplot(df.yrsurv[df.discent==1], hue=df['Age-Group'][df.discent==1], palette="Reds_r");

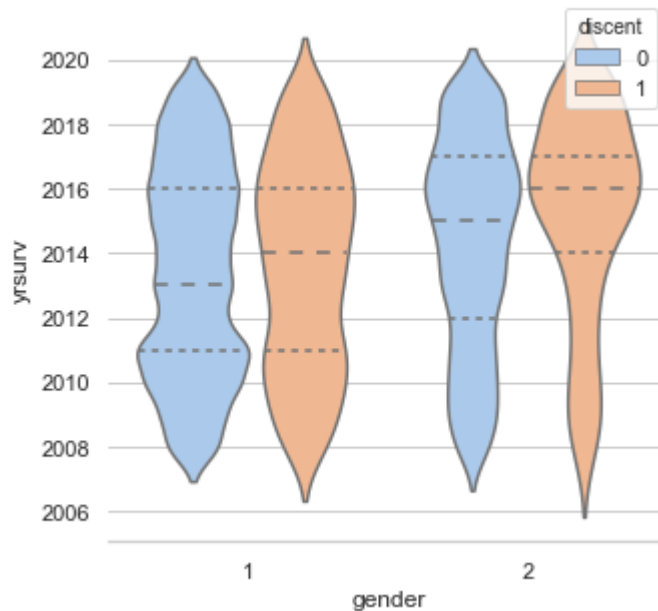
#sns.countplot(df1.yrsurv[df1.discent==0], hue=df['Age-Group'][df1.discent==0], palette="Blues_r");
plt.ylabel('Count of Exit')
plt.title('Number of Exit in Age group')
plt.show()
```



```
In [122]: sns.set(style="whitegrid", palette="pastel")


# Draw a nested violinplot and split the violins for easier comparison
sns.violinplot(x="gender", y="ysurv", hue='discent',
               split=False, inner="quart", data=df)

sns.despine(left=True)
```



```
In [123]: fig = plt.subplots(figsize=(10,5))

sns.countplot(df['Age-Group'], hue=df['gender'], palette="Blues_r");
#sns.countplot(df['Age-Group'], hue=df['discent'], palette="Reds_r");
```



```
In [157]: print(df.discent.value_counts())
```


```
sns.countplot(df.discent)
```

```
0    11022
```

```
1     1078
```

```
Name: discent, dtype: int64
```

```
Out[157]: <matplotlib.axes._subplots.AxesSubplot at 0x14c88a104c8>
```




```
In [159]: from pylab import rcParams

# Data to plot
sizes = df['discent'].value_counts(sort = True)
colors = ["grey", "yellow"]
rcParams['figure.figsize'] = 5,5 # Plot

plt.pie(sizes, explode = (0,0.3), labels=["NO", "YES"], colors=colors,
        autopct='%.2f%%', shadow=True, startangle=270)

plt.title('Percentage of EXIT in Business')
```

Out[159]: Text(0.5, 1.0, 'Percentage of EXIT in Business')



```
In [160]: df=df.drop(['Age-Group'],axis=1)
```

In [161]: df

Out[161]:

	yrsurv	age	gender	irstrata	fearfail	suskill	opport	knowent	discent	exreason	bstar
0	2008	28	2	2	2	1	1	1	0	0	1
1	2008	38	1	2	2	1	1	1	0	0	1
2	2008	26	2	6	2	1	1	1	0	0	1
3	2008	25	1	6	2	1	1	2	0	0	2
4	2008	45	1	2	-1	2	-1	1	0	0	1
...	...	...	...	...	...	...	...	...	...	...	..
12095	2019	54	2	6	1	1	2	1	0	0	1
12096	2019	21	2	6	1	2	2	1	0	0	1
12097	2019	24	2	6	2	1	1	1	0	0	1
12098	2019	31	2	6	2	1	2	2	0	0	1
12099	2019	32	2	6	2	2	2	1	0	0	2

12100 rows × 11 columns

In [162]: x=df.drop(['discent'], axis=1)  
y=df.discent

In [163]: print(x.shape , y.shape)  
  
(12100, 10) (12100,)

In [164]: x\_train, x\_test, y\_train, y\_test= train\_test\_split(x, y, test\_size=0.3, random\_state=1)  
  
x\_train, x\_val, y\_train, y\_val = train\_test\_split(x\_train, y\_train, train\_size=0.7, random\_state=1)  
# 0.25 x 0.8 = 0.2  
  
print('valid', x\_val.shape, y\_val.shape)  
print('Train', x\_train.shape, y\_train.shape)  
print('Test', x\_test.shape, y\_test.shape)  
  
valid (2541, 10) (2541,)  
Train (5929, 10) (5929,)  
Test (3630, 10) (3630,)

In [165]: from imblearn import over\_sampling  
from imblearn.over\_sampling import SMOTE  
import sklearn  
import imblearn

```
In [166]: y_train.value_counts()
```

```
Out[166]: 0    5394  
          1     535  
          Name: discent, dtype: int64
```


```
In [167]: smt = SMOTE()  
x_train , y_train = smt.fit_sample(x_train , y_train)
```

```
In [168]: np.bincount(y_train)
```

```
Out[168]: array([5394, 5394], dtype=int64)
```

```
In [169]: sns.countplot(y_train)
```

```
Out[169]: <matplotlib.axes._subplots.AxesSubplot at 0x14c8886f748>
```



```
In [170]: x_train.shape
```

```
Out[170]: (10788, 10)
```

```
In [171]: y_train.value_counts()
```

```
Out[171]: 1    5394  
          0    5394  
          Name: discent, dtype: int64
```



```
In [210]: from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC

from xgboost import XGBClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score , average_precision_score , precision_recall_curve , f1_score , roc_curve , auc
```

```
In [211]: #LinearRegression
lm = LinearRegression()

lm.fit(x_train, y_train)

y_pred = lm.predict(x_test)

lm.score(x_test, y_test)
```

Out[211]: 0.5497208065285251

```
In [242]: #LogisticRegression
lr_c=LogisticRegression(random_state=0)
lr_c.fit(x_train,y_train)
lr_pred=lr_c.predict(x_val)
lr_cm=confusion_matrix(y_val,lr_pred)
lr_ac=accuracy_score(y_val, lr_pred)
```

```
In [243]: #RandomForest
rdf_c=RandomForestClassifier(n_estimators=10,criterion='entropy',random_state=0)
rdf_c.fit(x_train,y_train)
rdf_pred=rdf_c.predict(x_val)
rdf_cm=confusion_matrix(y_val,rdf_pred)
rdf_ac=accuracy_score(rdf_pred,y_val)
```

```
In [244]: # DecisionTree Classifier
dtree_c=DecisionTreeClassifier(criterion='entropy',random_state=0)
dtree_c.fit(x_train,y_train)
dtree_pred=dtree_c.predict(x_val)
dtree_cm=confusion_matrix(y_val,dtree_pred)
dtree_ac=accuracy_score(dtree_pred,y_val)
```

```
In [245]: #XGBoost
XGB=XGBClassifier()
XGB.fit(x_train,y_train)
XGB_pred=XGB.predict(x_val)
XGB_cm=confusion_matrix(y_val,XGB_pred)
XGB_ac=accuracy_score(XGB_pred,y_val)
```

```
In [257]: #SVM
svm_c=SVC( C=1.0)
svm_c.fit(x_train,y_train)
svm_pred=svm_c.predict(x_val)
svm_cm=confusion_matrix(y_val,svm_pred)
svm_ac=accuracy_score(svm_pred,y_val)
```

```
In [ ]:
```

```

In [258]: plt.figure(figsize=(20,10))
target_names = ["continue", "exit"]

plt.subplot(2,3,1)
plt.title("LogisticRegression_cm")
sns.heatmap(lr_cm,annot=True,cmap="Paired_r",fmt="d",cbar=False)

plt.subplot(2,3,2)
plt.title("RandomForest")
sns.heatmap(rdf_cm,annot=True,cmap="BuPu",fmt="d",cbar=False)

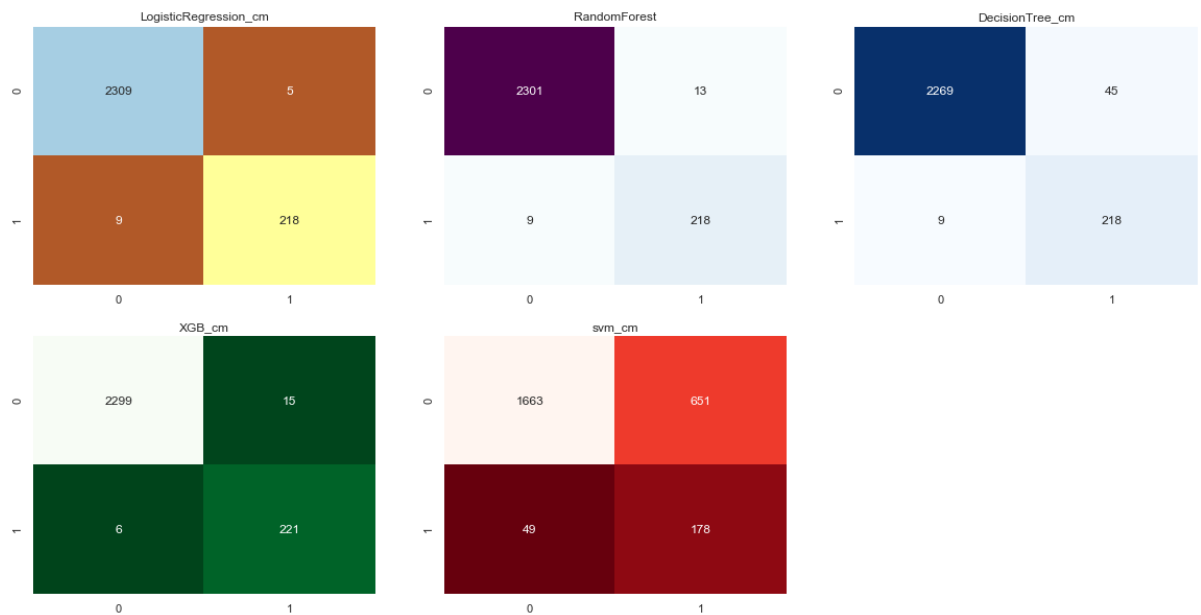
plt.subplot(2,3,3)
plt.title("DecisionTree_cm")
sns.heatmap(dtree_cm,annot=True,cmap="Blues",fmt="d",cbar=False)

plt.subplot(2,3,4)
plt.title("XGB_cm")
sns.heatmap(XGB_cm,annot=True,cmap="Greens_r",fmt="d",cbar=False)

plt.subplot(2,3,5)
plt.title("svm_cm")
sns.heatmap(svm_cm,annot=True,cmap="Reds_r",fmt="d",cbar=False)

```

Out[258]: <matplotlib.axes.\_subplots.AxesSubplot at 0x14c88921b48>



```
In [266]: print('LogisticRegression_accuracy:\t',lr_ac)
print('RandomForest_accuracy:\t\t',rdf_ac)
print('DecisionTree_accuracy:\t\t',dtree_ac)
print('XGB_accuracy:\t\t\t',XGB_ac)
print('SVM_accuracy:\t\t\t',svm_ac)
```

```
LogisticRegression_accuracy:    0.9944903581267218
RandomForest_accuracy:          0.9913419913419913
DecisionTree_accuracy:          0.9787485242030697
XGB_accuracy:                   0.9917355371900827
SVM_accuracy:                   0.7245179063360881
```

```
In [260]: def plotting(true,pred):
    fig,ax=plt.subplots(1,2,figsize=(10,5))


    precision,recall,threshold = precision_recall_curve(true,pred[:,1])
    ax[0].plot(recall,precision,'g--')
    ax[0].set_xlabel('Recall')
    ax[0].set_ylabel('Precision')
    ax[0].set_title("Average Precision Score : {}".format(average_precision_score(true,pred[:,1])))

    fpr,tpr,threshold = roc_curve(true,pred[:,1])
    ax[1].plot(fpr,tpr)
    ax[1].set_title("AUC Score is: {}".format(auc(fpr,tpr)))
    ax[1].plot([0,1],[0,1],'k--')
    ax[1].set_xlabel('False Positive Rate')
    ax[1].set_ylabel('True Positive Rate')
```

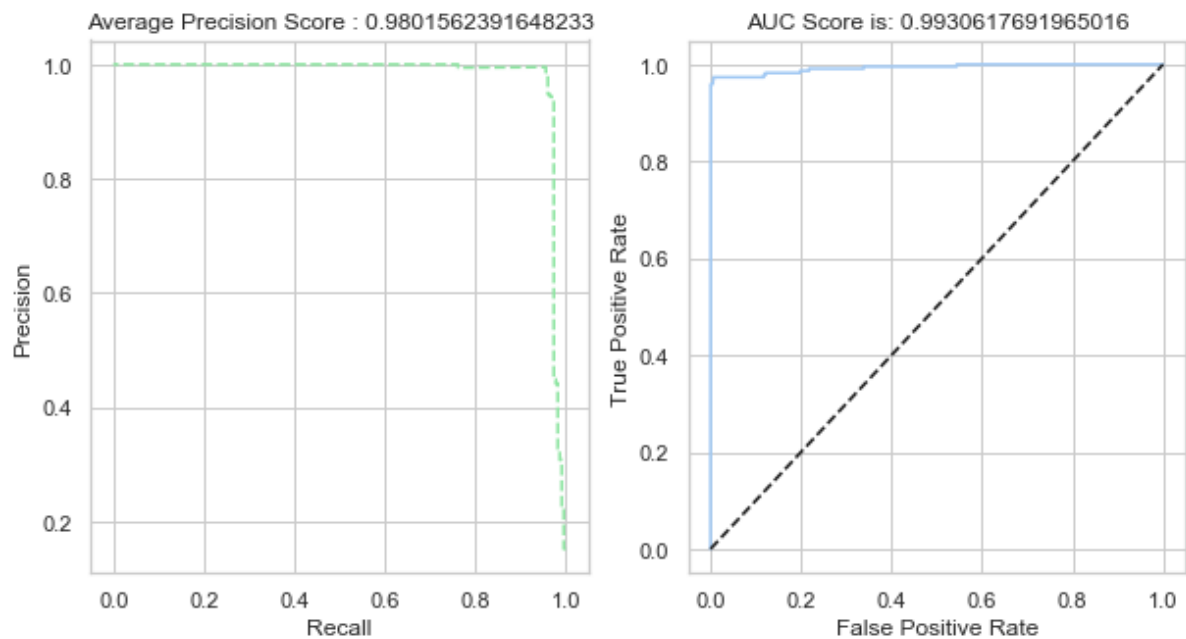
```
In [261]: plt.figure()
plotting(y_val ,rdf_c.predict_proba(x_val))

plt.figure()
plotting(y_val ,XGB.predict_proba(x_val))
```

<Figure size 360x360 with 0 Axes>



<Figure size 360x360 with 0 Axes>



```
In [262]: from sklearn.metrics import accuracy_score, confusion_matrix, precision_score,
recall_score, precision_recall_curve
```

```
print('precision_score of LR: ', precision_score(lr_pred, y_val))
print('precision_score of random forest : ', precision_score(rdf_pred, y_val
))
print('precision_score of DTree: ', precision_score(dtrees_pred, y_val))
print('precision_score of XGB: ', precision_score(XGB_pred, y_val))
print('precision_score of SVM: ', precision_score(svm_pred, y_val))
```

```
precision_score of LR: 0.960352422907489
precision_score of random forest : 0.960352422907489
precision_score of DTree: 0.960352422907489
precision_score of XGB: 0.973568281938326
precision_score of SVM: 0.7841409691629956
```

```
In [263]: print('recall_score of LR: ', recall_score(lr_pred, y_val))
print('recall_score of random forest : ', recall_score(rdf_pred, y_val))
print('recall_score of DTree: ', recall_score(dtrees_pred, y_val))
print('recall_score of XGB: ', recall_score(XGB_pred, y_val))
print('recall_score of SVM: ', recall_score(svm_pred, y_val))
```

```
recall_score of LR: 0.9775784753363229
recall_score of random forest : 0.9437229437229437
recall_score of DTree: 0.8288973384030418
recall_score of XGB: 0.9364406779661016
recall_score of SVM: 0.2147165259348613
```

```
In [264]: print('f1_score of LR: ', f1_score(lr_pred, y_val))
print('f1_score of random forest : ', f1_score(rdf_pred, y_val))
print('f1_score of DTree: ', f1_score(dtrees_pred, y_val))
print('f1_score of XGB: ', f1_score(XGB_pred, y_val))
print('f1_score of SVM: ', f1_score(svm_pred, y_val))
```

```
f1_score of LR: 0.9688888888888889
f1_score of random forest : 0.9519650655021834
f1_score of DTree: 0.889795918367347
f1_score of XGB: 0.9546436285097192
f1_score of SVM: 0.33712121212121215
```


```
In [265]: from sklearn import model_selection
dataframe = df
array = dataframe.values

# prepare configuration for cross validation test harness
seed = 7
# prepare models
models = []
models.append(('LR', LogisticRegression()))
models.append(('random forest', RandomForestClassifier(n_estimators=10, criterion='entropy', random_state=0)))
models.append(('DTree', DecisionTreeClassifier(criterion='entropy', random_state=0)))
models.append(('XGB', XGBClassifier()))
models.append(('svm', SVC()))

# evaluate each model in turn
results = []
names = []
scoring = 'accuracy'
for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state=seed)
    cv_results = model_selection.cross_val_score(model, x, y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
# boxplot algorithm comparison
fig = plt.figure()
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

LR: 0.993140 (0.004527)  
 random forest: 0.993554 (0.006270)  
 DTree: 0.982893 (0.018947)  
 XGB: 0.993802 (0.006176)  
 svm: 0.910909 (0.010653)

Algorithm Comparison



```
In [225]: from sklearn.model_selection import GridSearchCV
rf = RandomForestClassifier()
n_estimators = [500,800,1000,2000]
max_depth = [3,5,8,10,11]
min_samples_split = [3,5]
min_samples_leaf = [1]

hyperF = dict (n_estimators = n_estimators, max_depth = max_depth,
               min_samples_split = min_samples_split,
               min_samples_leaf = min_samples_leaf)

gridF = GridSearchCV(rf , hyperF, cv = 3, verbose = 1)
bestF = gridF.fit(x_train, y_train)
```

Fitting 3 folds for each of 40 candidates, totalling 120 fits

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n\_jobs=1)]: Done 120 out of 120 | elapsed: 17.7min finished

```
In [226]: gridF.best_params_
```

```
Out[226]: {'max_depth': 11,
           'min_samples_leaf': 1,
           'min_samples_split': 3,
           'n_estimators': 800}
```



```
In [227]: gridF.best_score_
```

```
Out[227]: 0.9839636633296255
```

```
In [229]: model = RandomForestClassifier(n_estimators = 800, max_depth = 11,
        min_samples_split = 3,
        min_samples_leaf = 1)
model.fit(x_val, y_val)

model_pred=model.predict(x_val)
model_cm=confusion_matrix(y_val,model_pred)
model_ac=accuracy_score(model_pred,y_val)
```

```
In [230]: print('RandomForest_accuracy:\t\t',model_ac)

print('recall_score : ', recall_score(model_pred, y_val))

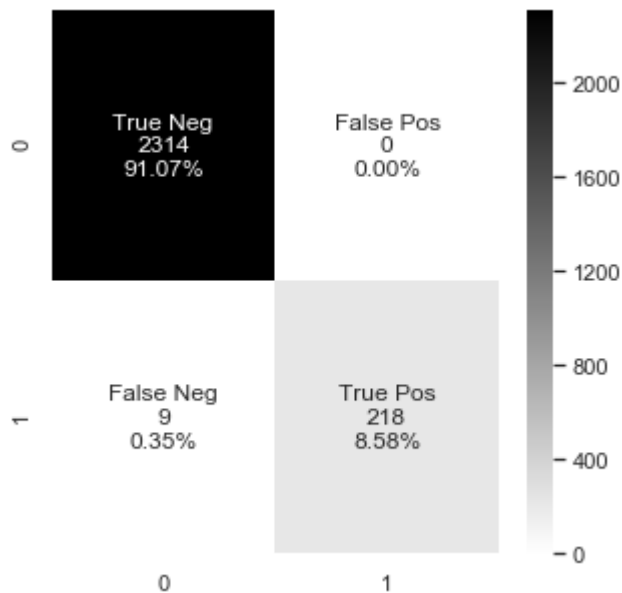
print('precision_score : ', precision_score(model_pred, y_val))

print('f1_score : ', f1_score(model_pred, y_val))
```

```
RandomForest_accuracy:          0.9964580873671782
recall_score : 1.0
precision_score : 0.960352422907489
f1_score : 0.9797752808988764
```

```
In [231]: group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
group_counts = ['{0:0.0f}'.format(value) for value in
                model_cm.flatten()]
group_percentages = ['{0:.2%}'.format(value) for value in
                    model_cm.flatten()/np.sum(model_cm)]
labels = [f"{v1}\n{v2}\n{v3}" for v1, v2, v3 in
          zip(group_names, group_counts, group_percentages)]
labels = np.asarray(labels).reshape(2,2)
sns.heatmap(model_cm, annot=labels, fmt='', cmap='binary')
```

Out[231]: <matplotlib.axes.\_subplots.AxesSubplot at 0x14c88e95d88>



```
In [232]: accuracy = model.score(x_test, y_test)
print("Accuracy is %.2f %" % (accuracy * 100))
```

Accuracy is 99.37 %

In [233]: `pred_classes = model.predict(x_test)`

```
print("Predicted classes:")
print(pred_classes)
print("Actual classes:")
print(y_test)
```

Predicted classes:

[0 0 1 ... 0 1 0]

Actual classes:

1300 0

8033 0

5853 1

1475 0

1416 0

..

11861 0

3933 0

3906 1

4760 1

11628 0

Name: discent, Length: 3630, dtype: int64

In [234]: `from sklearn.metrics import roc_auc_score`

In [235]: `y_pred = model.predict(x_test)`  
`print("Roc_auc_score: ",roc_auc_score(y_test,y_pred)*100,"%")`

Roc\_auc\_score: 96.36075949367088 %

In [236]: `model_T = RandomForestClassifier(n_estimators = 800, max_depth = 11,`  
`min_samples_split = 3,`  
`min_samples_leaf = 1)`  
`model_T.fit(x_test, y_test)`  
  
`model_pred=model_T.predict(x_test)`  
`model_cm=confusion_matrix(y_test,model_pred)`  
`model_ac=accuracy_score(model_pred,y_test)`

In [237]: `from sklearn.metrics import accuracy_score, confusion_matrix, precision_score,`  
`recall_score, precision_recall_curve`  
`print('RandomForest_accuracy:\t\t',model_ac)`  
  
`print('recall_score : ', recall_score(model_pred, y_test))`  
  
`print('precision_score : ', precision_score(model_pred, y_test))`  
  
`print('f1_score : ', f1_score(model_pred, y_test))`

RandomForest\_accuracy: 0.9939393939393939

recall\_score : 1.0

precision\_score : 0.930379746835443

f1\_score : 0.9639344262295081

In [ ]:

In [ ]:


```
In [238]: def plotting(true,pred):
            fig,ax=plt.subplots(1,2,figsize=(10,5))

            precision,recall,threshold = precision_recall_curve(true,pred[:,1])
            ax[0].plot(recall,precision,'g--')
            ax[0].set_xlabel('Recall')
            ax[0].set_ylabel('Precision')
            ax[0].set_title("Average Precision Score : {}".format(average_precision_score(true,pred[:,1])))

            fpr,tpr,threshold = roc_curve(true,pred[:,1])
            ax[1].plot(fpr,tpr)
            ax[1].set_title("AUC Score is: {}".format(auc(fpr,tpr)))
            ax[1].plot([0,1],[0,1],'k--')
            ax[1].set_xlabel('False Positive Rate')
            ax[1].set_ylabel('True Positive Rate')
```

```
In [239]: plt.figure()
            plotting(y_val ,model.predict_proba(x_val))
```

<Figure size 360x360 with 0 Axes>




```
In [240]: import sklearn.metrics as metrics
def plotting(true,pred):
    fig,ax=plt.subplots(1,2,figsize=(10,5))

    precision,recall,threshold = precision_recall_curve(true,pred[:,1])
    ax[0].plot(recall,precision,'g--')
    ax[0].set_xlabel('Recall')
    ax[0].set_ylabel('Precision')
    ax[0].set_title("Average Precision Score : {}".format(average_precision_score(true,pred[:,1])))

    fpr,tpr,threshold = roc_curve(true,pred[:,1])
    roc_auc = metrics.auc(fpr, tpr)
    ax[1].plot(fpr,tpr)
    ax[1].set_title("AUC Score is: {}".format(auc(fpr,tpr)))
    ax[1].plot([0,1],[0,1], 'k--')
    ax[1].set_xlabel('False Positive Rate')
    ax[1].set_ylabel('True Positive Rate')
```

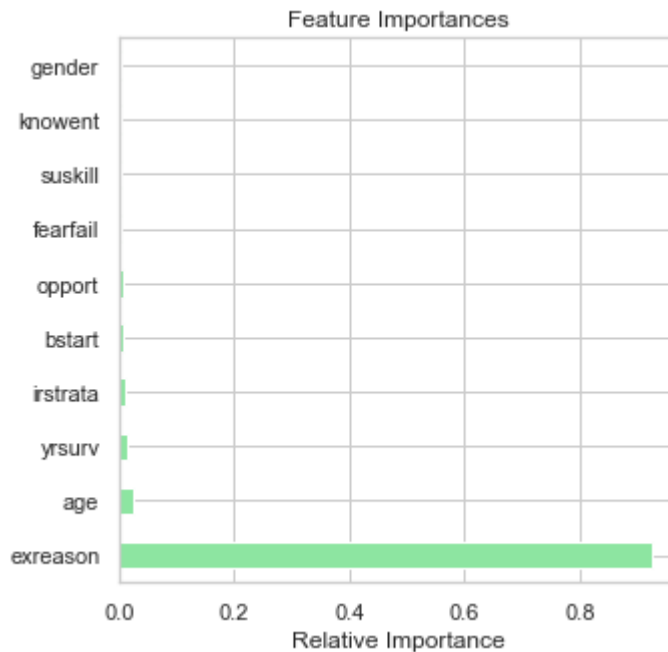
```
In [241]: plt.figure()
plotting(y_test ,model.predict_proba(x_test))
```

<Figure size 360x360 with 0 Axes>



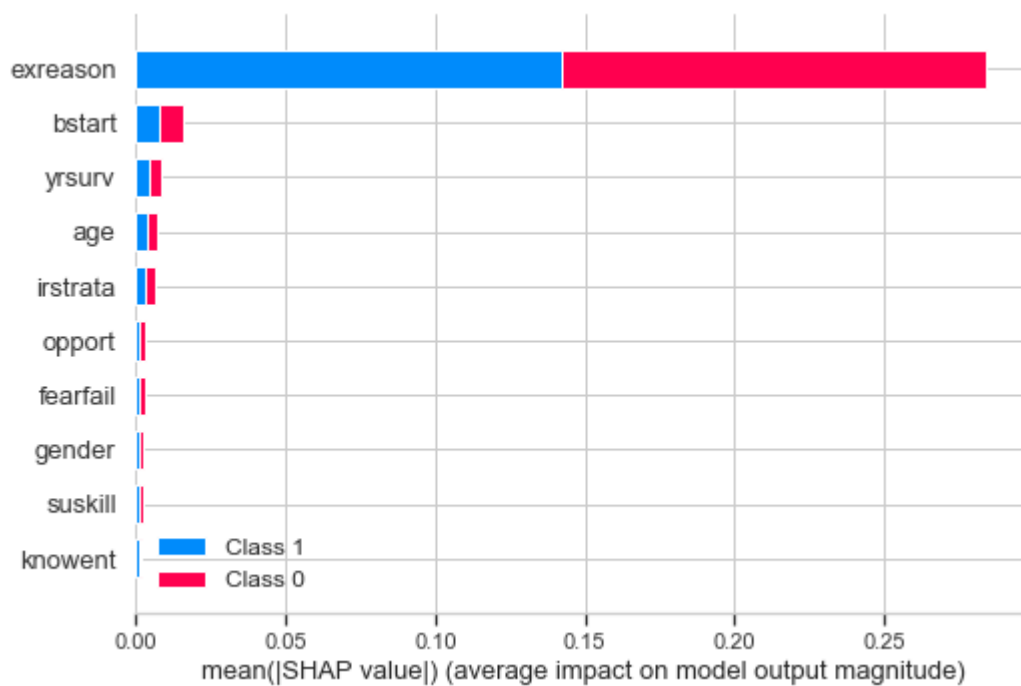
```
In [109]: feat_importances = pd.Series(model.feature_importances_, index=x.columns)

feat_importances.nlargest(20).plot(kind='barh', color='g',)
plt.title('Feature Importances')
#plt.barh(range(len(indices)), importances[indices], color='g', align='center')
#plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



```
In [92]: import shap
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(x_test)
```

```
In [93]: shap.summary_plot(shap_values, x_test, plot_type="bar")
```



```
In [94]: from sklearn.model_selection import learning_curve  
from sklearn.model_selection import ShuffleSplit
```

```

In [95]: train_sizes, train_scores, test_scores = learning_curve(RandomForestClassifier(
    (),
    x,
    y,
    # Number of folds in cross-validation
    cv=10,
    # Evaluation metric
    scoring='accuracy',
    # Use all computer cores
    n_jobs=-1,
    # 50 different sizes of the training set
    train_sizes=np.linspace(0.01, 1.0, 50))

# Create means and standard deviations of training set scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)

# Create means and standard deviations of test set scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)


# Draw lines
plt.plot(train_sizes, train_mean, '--', color="#111111", label="Training score")
plt.plot(train_sizes, test_mean, color="#111111", label="Cross-validation score")

# Draw bands
plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std, color="#DDDDDD")
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std, color="#DDDDDD")

# Create plot
plt.title("Learning Curve")
plt.xlabel("Training Set Size"), plt.ylabel("Accuracy Score"), plt.legend(loc="best")
plt.tight_layout()
plt.show()

```





```
In [96]: # save the model to disk with pickle
from pickle import dump

filename = 'finalized_RandomForest_GEM_model-E.sav'
dump(model, open(filename, 'wb'))
```

```
In [97]: # Load the model from disk with pickle
from pickle import load

loaded_model = load(open(filename, 'rb'))
result = loaded_model.score(x_test, y_test)
print(result)
```

```
0.99366391184573
```

```
In [ ]:
```

```
In [ ]:
```